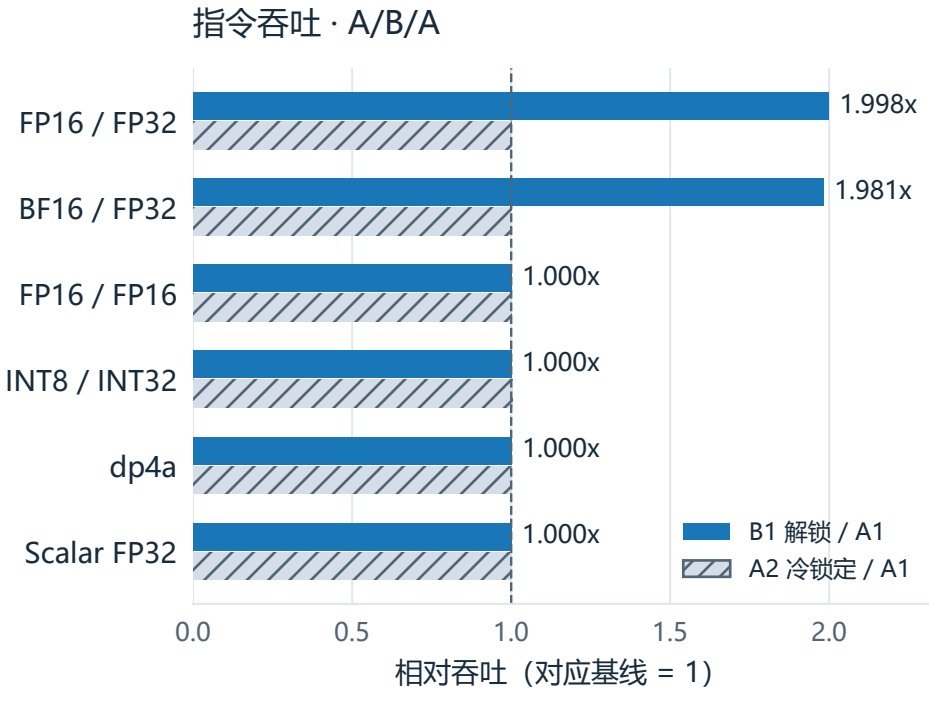
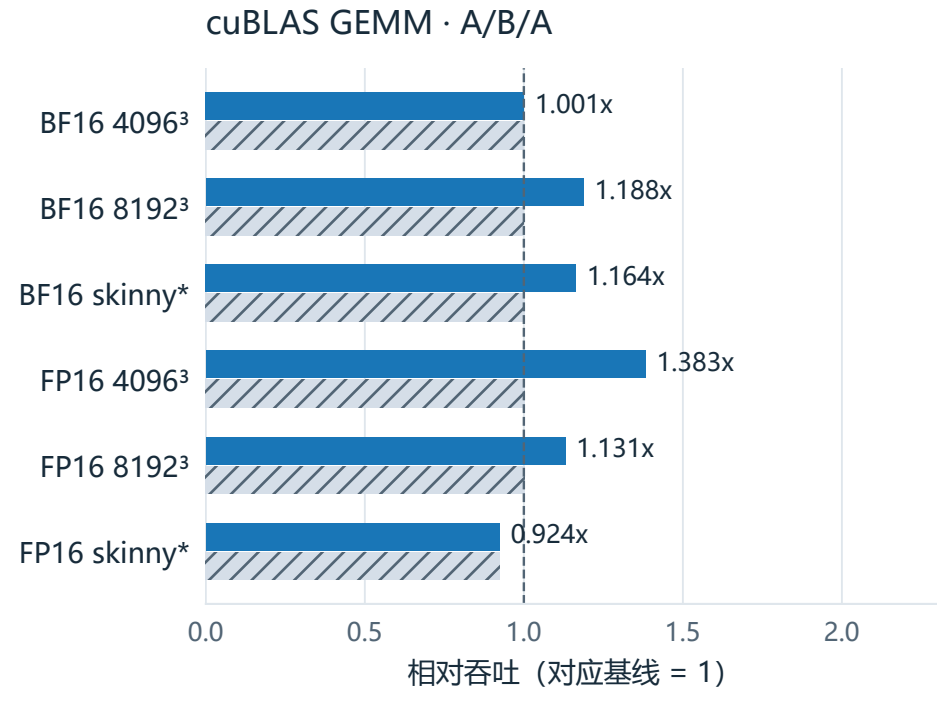
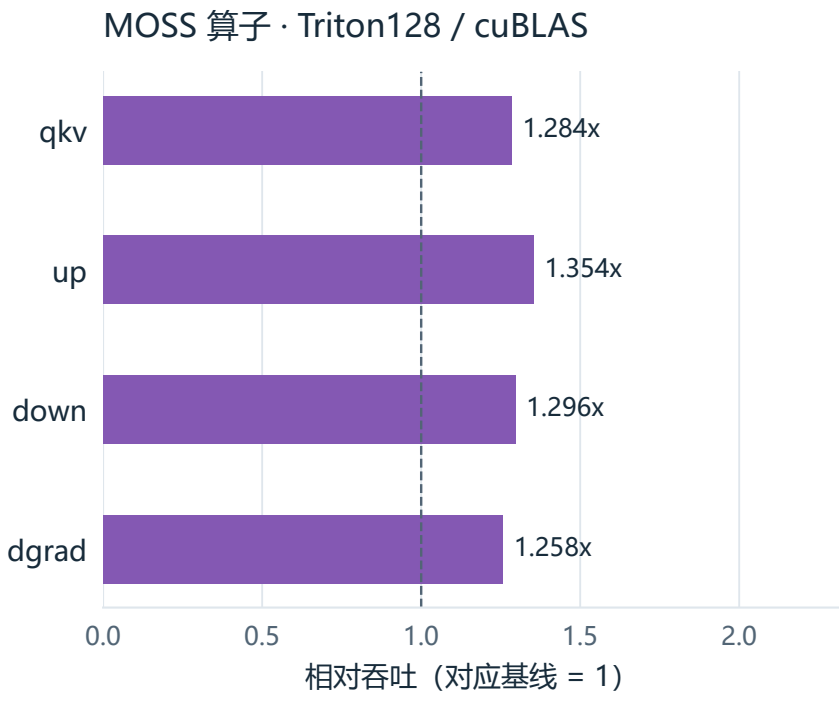
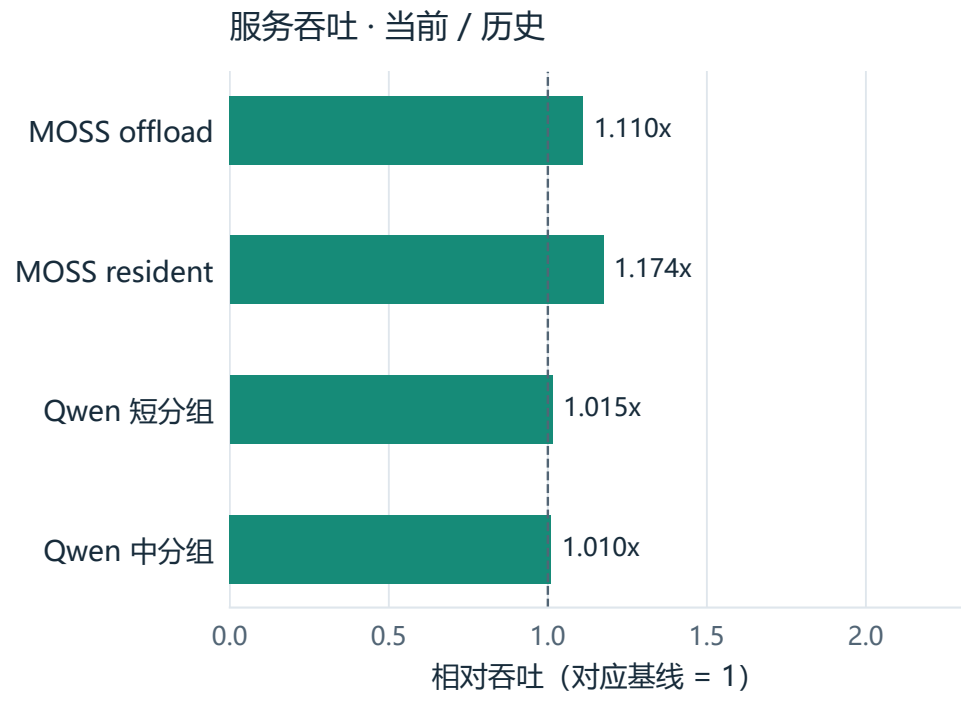
算力、算子与生产服务的性能对比

双 RTX 4090 · 48 GiB/卡 · 450 W/卡









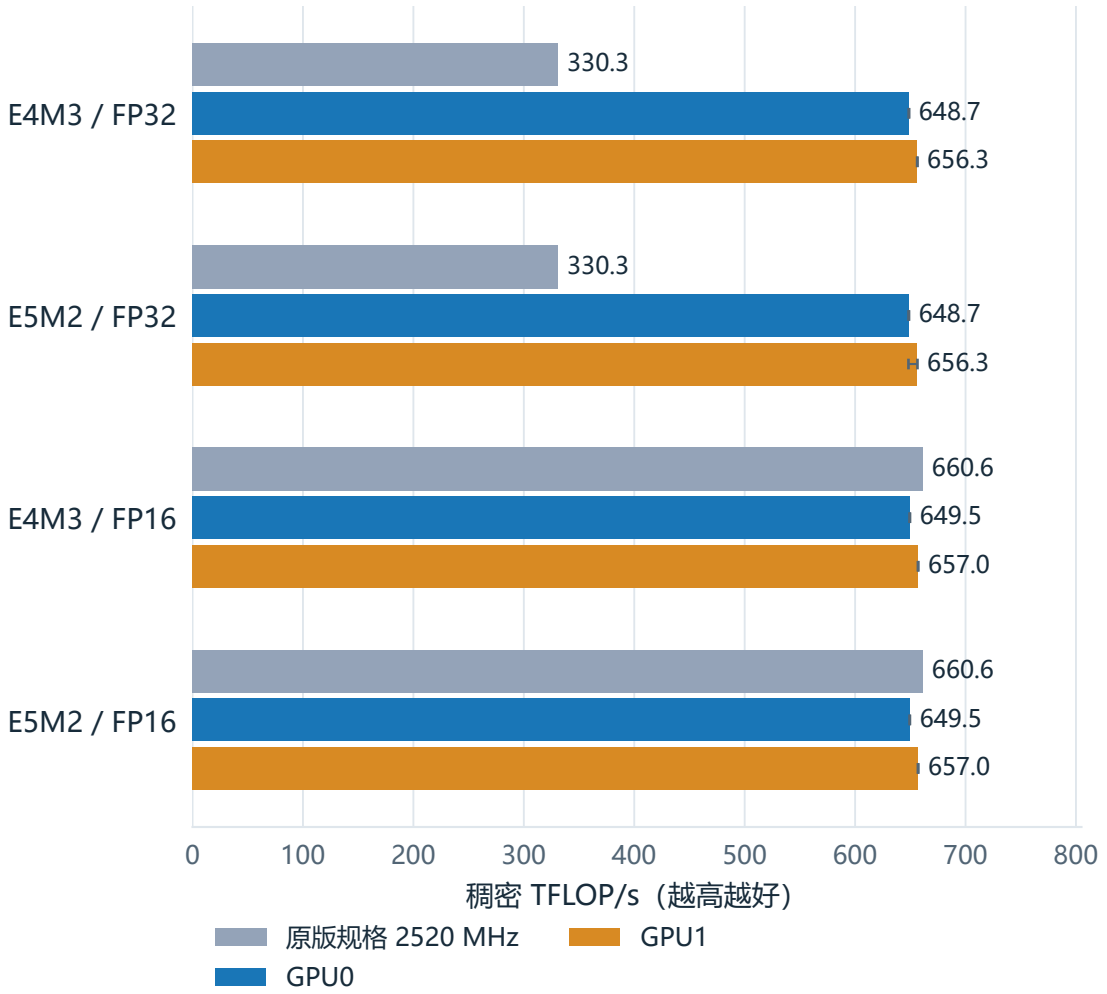
* skinny = 256×512×16384；指令 n=30/模式，GEMM n=10/形状。

Sources: ab-a-matched.json | operators.json | historical-applications.json (2026-10-26)

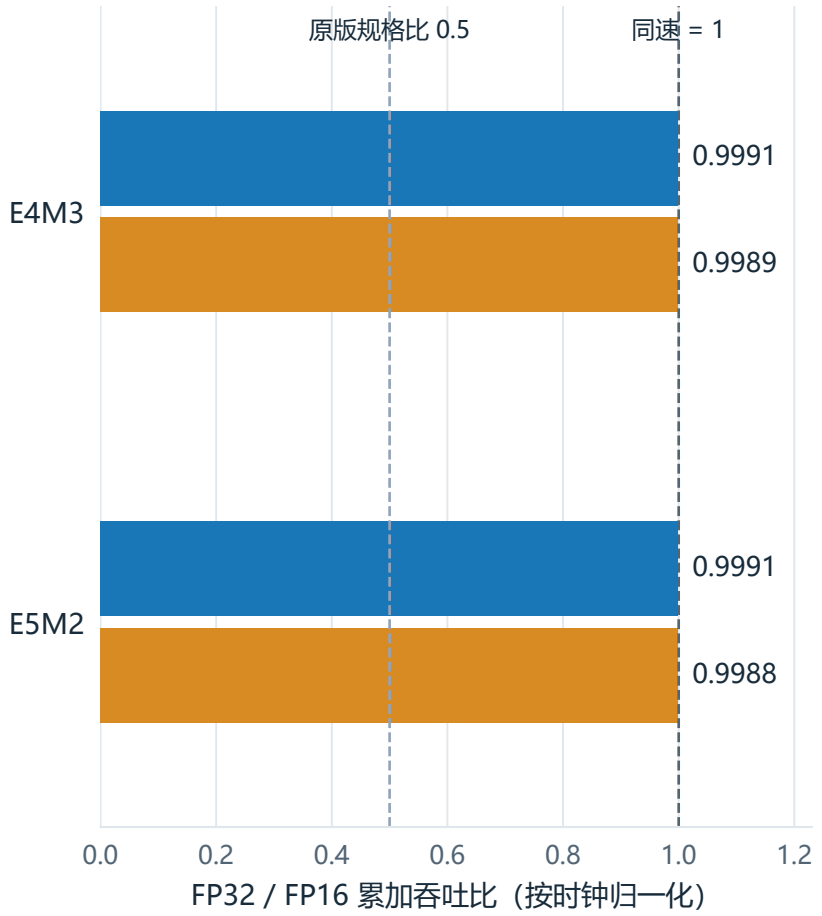
FP8 原生算力

已解锁双卡 · 每模式 n=21 · 稠密计算 · 误差线 min-max

原版规格参考与当前实测



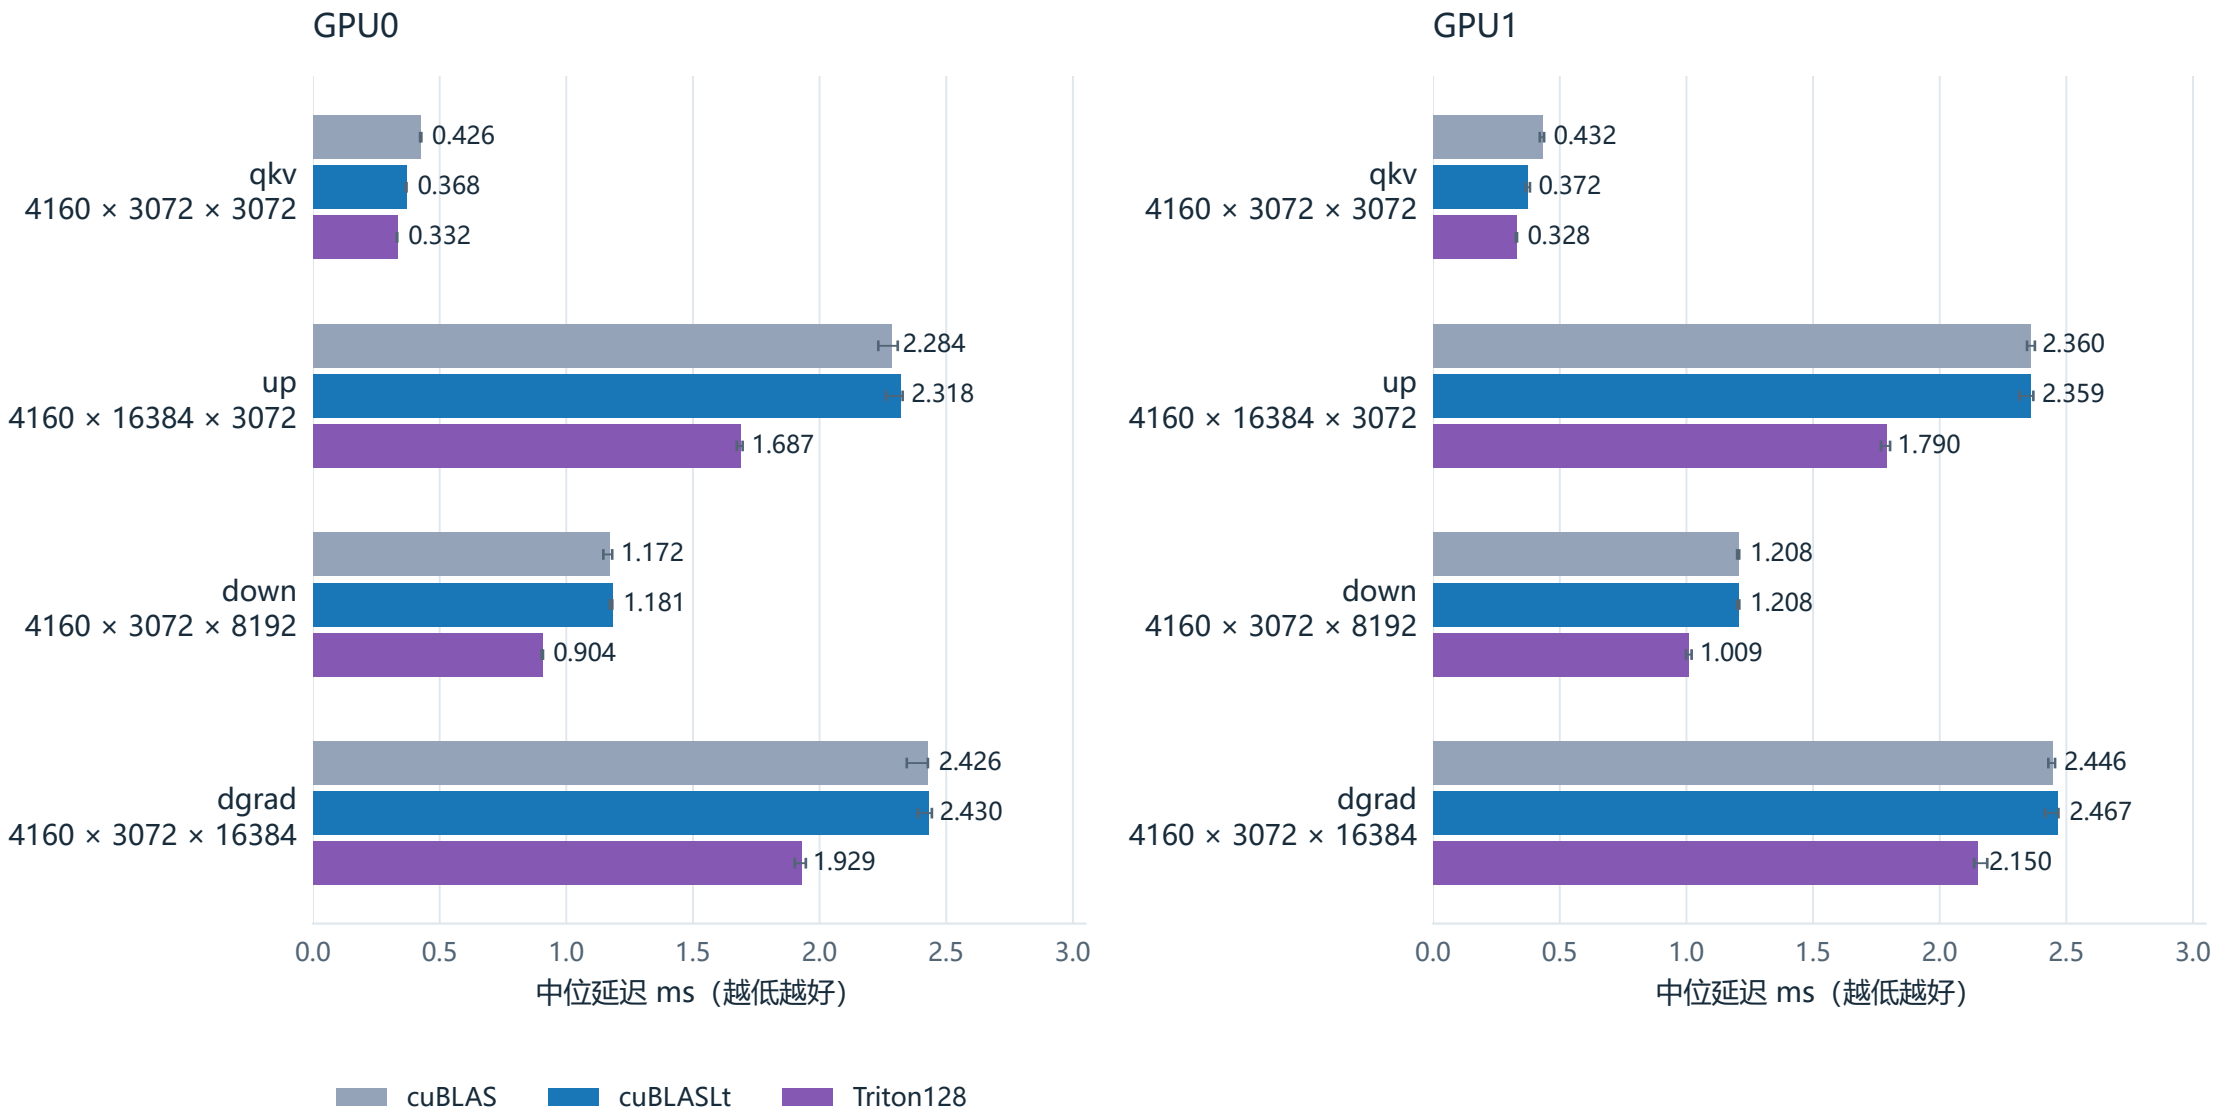
同卡 FP32 与 FP16 累加对照



时钟：GPU0 约 2710 MHz · GPU1 约 2744 MHz · 原版规格 2520 MHz。

MOSS 算子延迟

同一已解锁状态、相同 BF16 输入、450 W；每实现 35 次，误差线为观测 P10-P90。

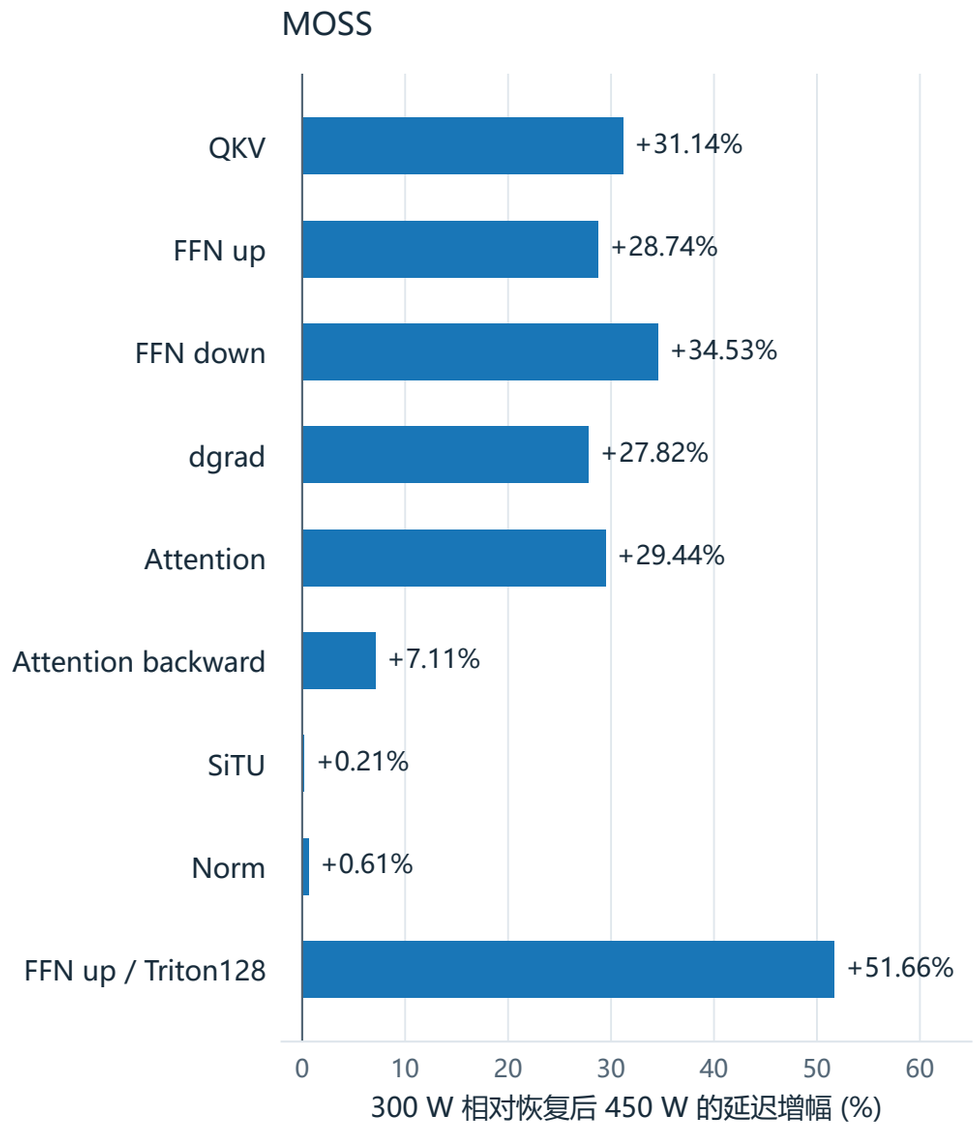


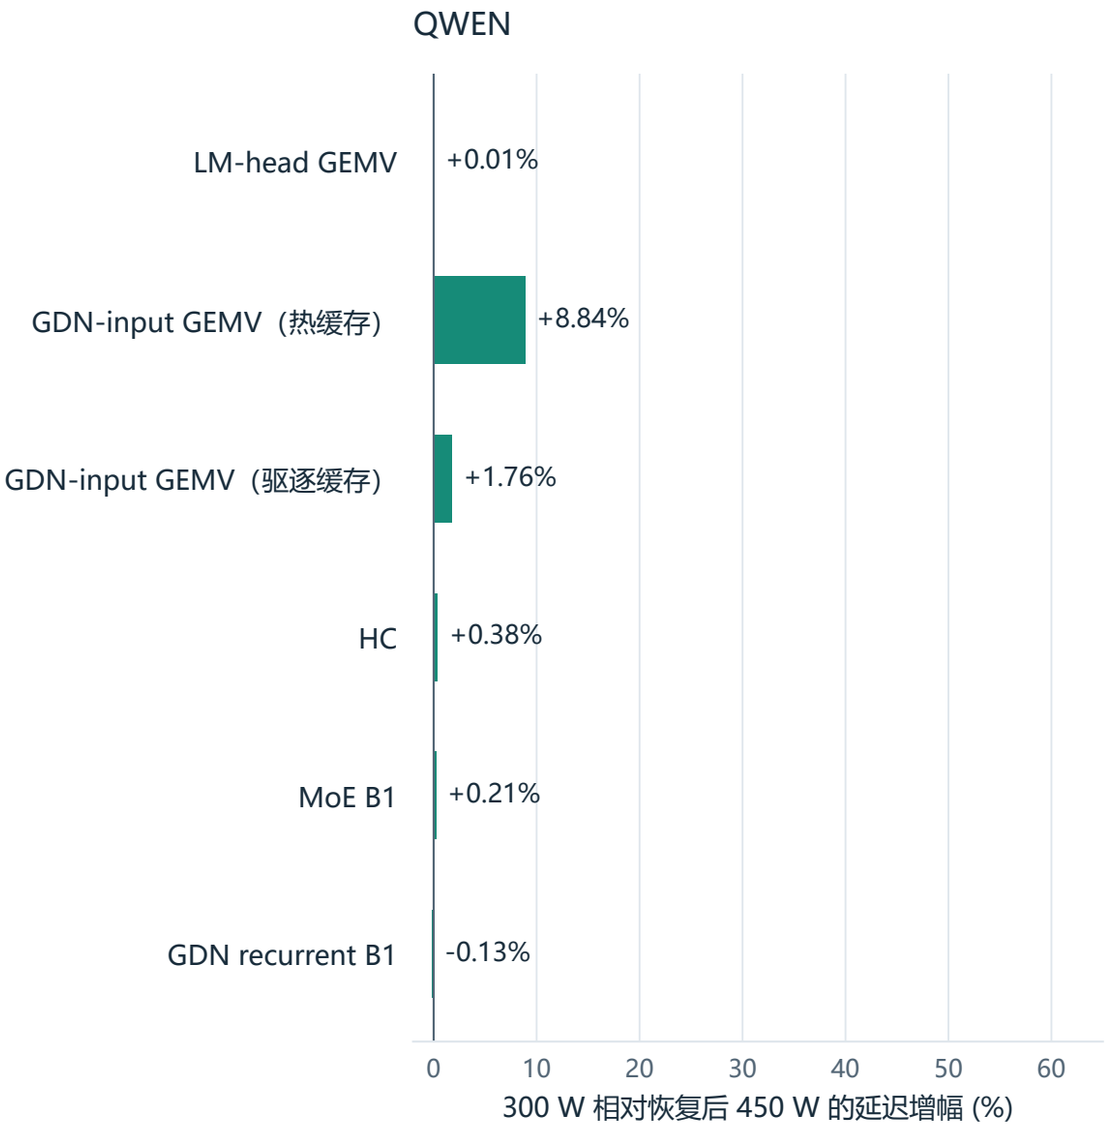
输入形状标为 M × N × K。

Source: 2026-10-26/operators.json, paired-g0-450 and paired-g1-450

MOSS / Qwen 算子功耗敏感度

同卡、同输入，比较 300 W 与恢复后的 450 W；两面板使用相同横轴。





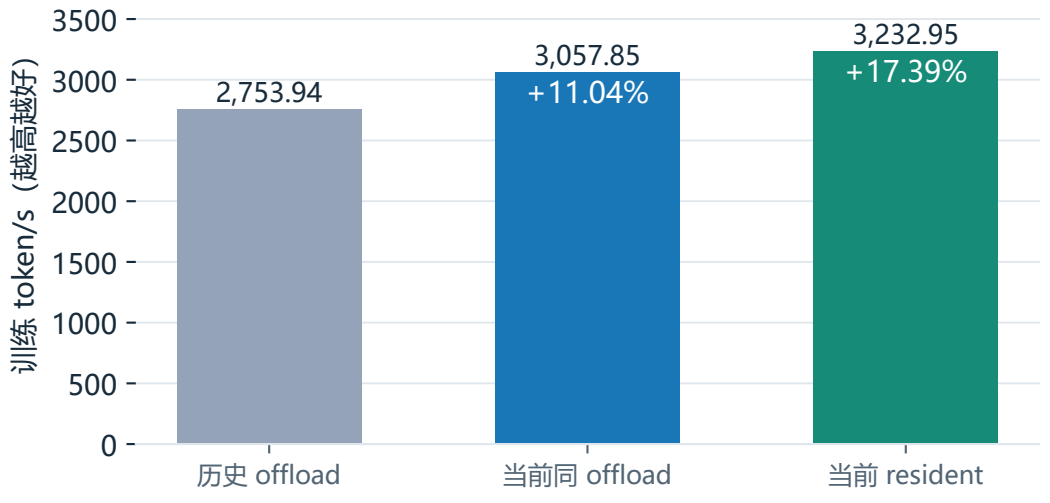
GPU0 · CUDA event 计时 · 常规 n=21，驱逐缓存 n=31。

Source: 2026-10-26/operators.json, *-g0-300 and *-g0-450b

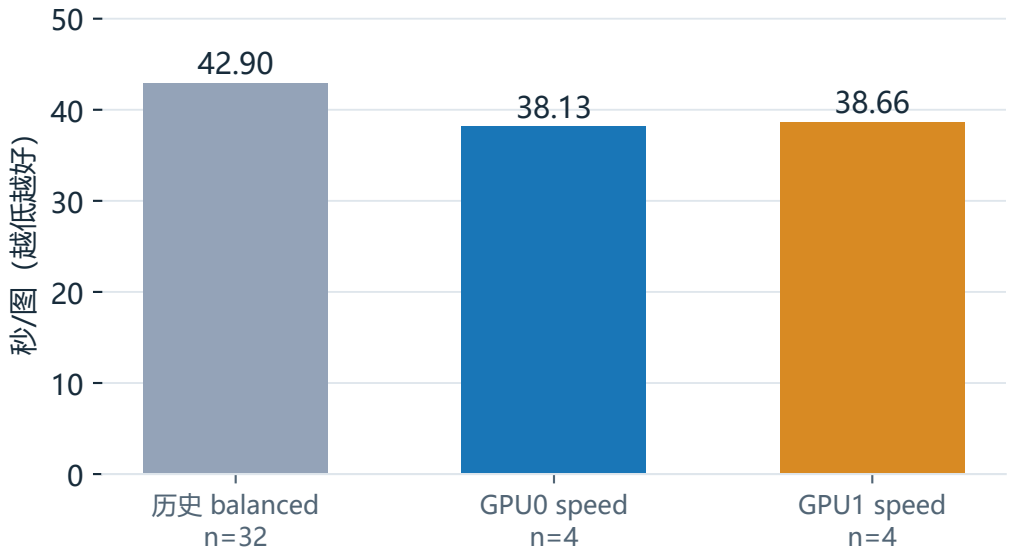
MOSS 训练与推理性能

训练吞吐 = 总 token / 总区间时间

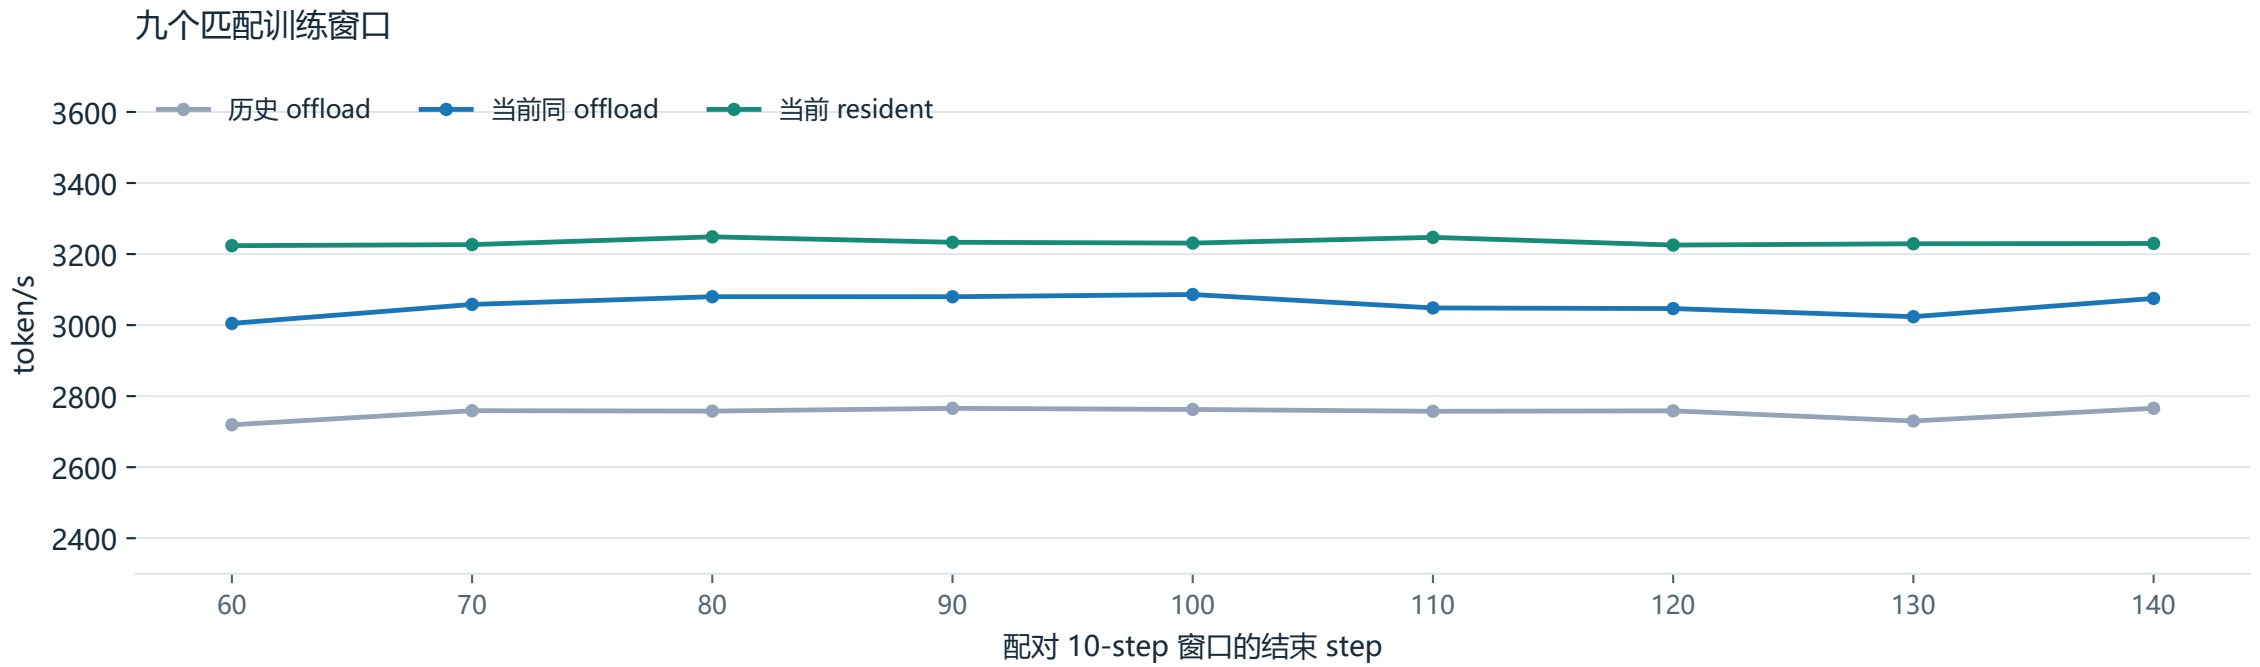
匹配训练区间：948,960 token / 90 step



图像生成



九个匹配训练窗口



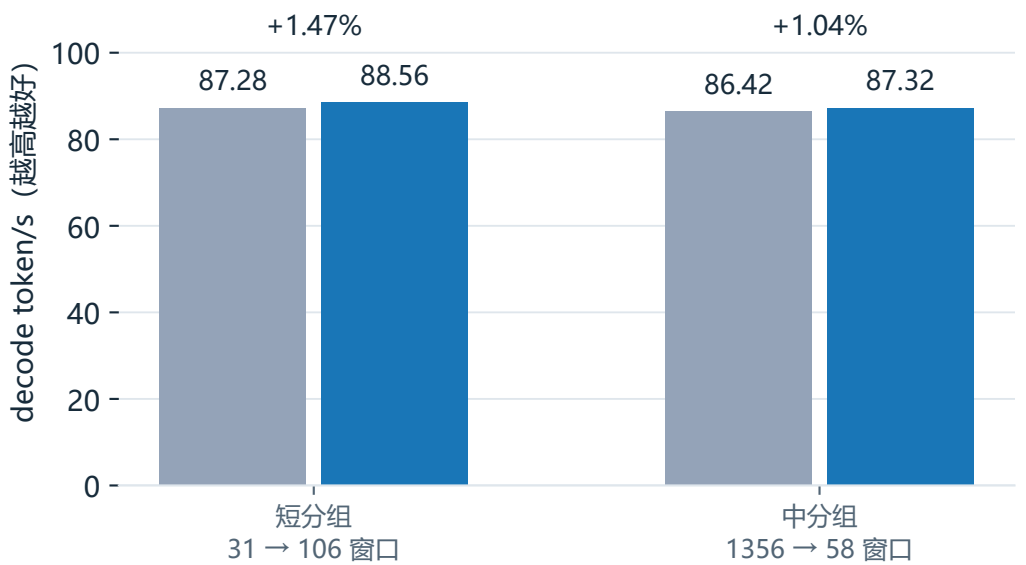
训练区间：step 51-140 · 90 step · 237 样本 · 948,960 token。

Source: 2026-10-26/historical-applications.json

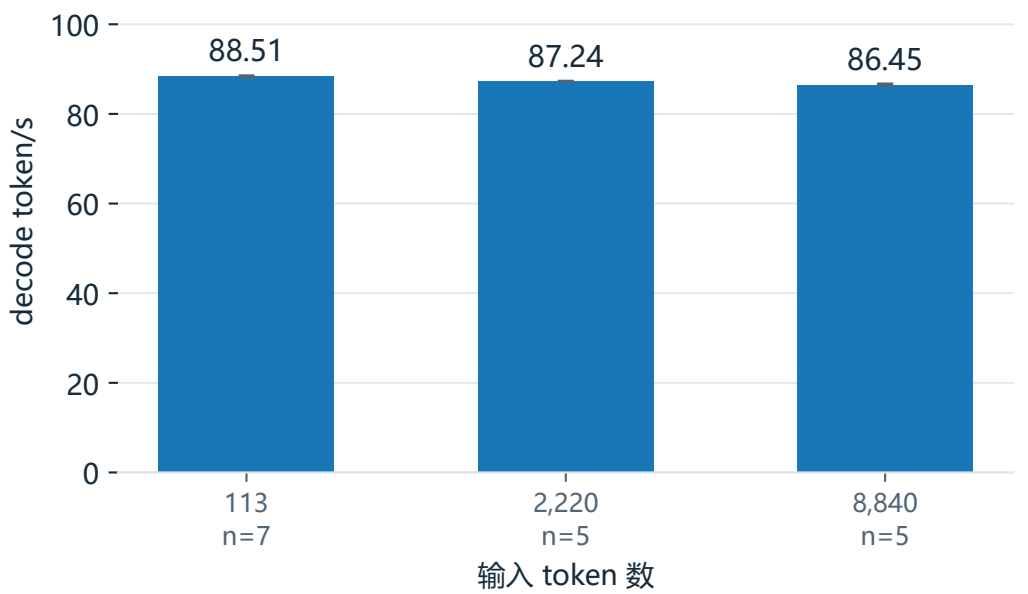
Qwen 服务性能

C1 · TP2 · prefix-cache hit = 0 · 每请求输出 512 token · 预热 2 次、测量 17 次

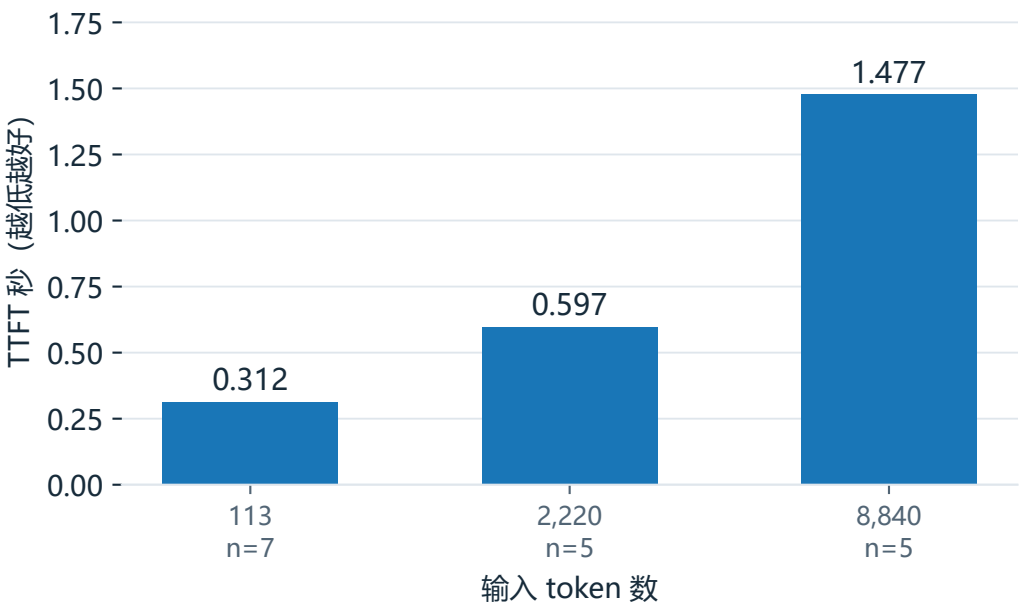
同 profile 历史日志



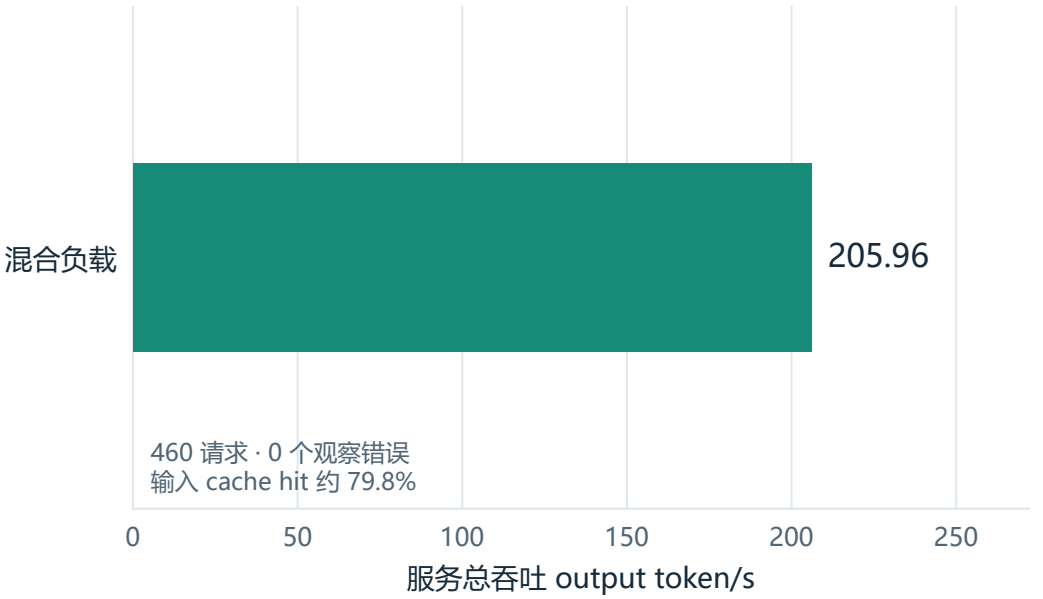
当前 C1：86-89 token/s



当前 C1：首 token 延迟



混合负载总吞吐



当前 C1 误差线为观测 min-max；历史日志样本数标在横轴。